

Stable and Convexified Information Bottleneck Optimization via Symbolic Continuation and Entropy-Regularized Trajectories

Faruk Alpay¹

¹Independent Researcher

¹ORCID: [0009-0009-2207-6528](https://orcid.org/0009-0009-2207-6528)

May 13, 2025

Abstract

The Information Bottleneck (IB) objective $I(X; Z) - \beta I(Z; Y)$ is known to exhibit instability and abrupt phase transitions as its trade-off parameter β is varied. These instabilities manifest as volatile jumps in the encoder $p(z|x)$ and sparse, hard cluster formations, posing challenges in high-sensitivity applications where smooth and stable representations are required. In this work, I introduce a convexified IB optimization framework that replaces the classical linear compression term with a strictly convex function $u(I(X; Z))$ (e.g., $u(t) = t^2$), and incorporate a small entropy regularization $-\epsilon H(Z|X)$ to smooth encoder transitions. To trace solutions continuously as β increases, I develop a symbolic continuation method based on an implicit first-order ODE for the encoder, which serves as a predictor-corrector mechanism for following the IB path without bifurcating. My approach yields a stable IB solver that avoids sudden representation shifts by design. I demonstrate on synthetic datasets (binary symmetric channel and structured 8×8 distributions) that the method I propose eliminates abrupt phase changes, achieving smooth evolution of mutual information metrics and graceful cluster formation. The convexified and entropy-regularized IB solutions maintain higher stability and predictive performance across all β regimes, confirming the theoretical improvements.

Keywords: Information Bottleneck; continuation methods; convex optimization; entropy regularization; phase transitions

1 Introduction

The Information Bottleneck (IB) framework [1] provides a principled approach for extracting a compressed representation Z of a source X that is maximally informative about a target Y . By optimizing the trade-off between compression $I(X; Z)$ and prediction quality $I(Z; Y)$, IB has shown success in applications ranging from robust deep learning representations [2] to invariant and disentangled feature learning [3]. Despite its appeal, a notorious issue in standard IB is the emergence of instabilities and phase transitions in the optimal encoder as the trade-off parameter β is varied. At certain critical β values, the IB solution undergoes abrupt changes – e.g. sudden jumps in $I(Z; Y)$ or discretization of $p(z|x)$ – corresponding to the spontaneous creation of new clusters in Z . Wu et al. [4] characterize these IB phase transitions as qualitative changes in the loss landscape that mark the onset of learning new classes or features. Such volatile behavior is problematic in high-sensitivity domains (e.g. medical decision systems or adaptive communication networks) where slight parameter perturbations should not cause discontinuous jumps in the learned representation. Ensuring a stable and smooth evolution of the IB representation is therefore of practical importance.

Several prior works have identified the root causes of IB instability. In deterministic scenarios where $Y = f(X)$, the IB curve (the achievable $I(Z; Y)$ vs. $I(X; Z)$ frontier) becomes piecewise linear and non-convex [8], leading to degenerate solutions where many β values yield the same optimum [6]. Kolchinsky et al. [8] proved that in such cases the standard IB Lagrangian cannot continuously explore intermediate compression levels [8]. Instead, the solution tends to jump from the trivial representation (Z independent of X) directly to a near-complete information retention, bypassing the intermediate regime. Even in non-deterministic settings, the IB objective often exhibits bifurcations: as β increases, a formerly optimal encoder can become unstable when a Hessian eigenvalue passes zero, at which point a new branch of solutions (with an extra cluster) forks off [5]. These instabilities are analogous to symmetry-breaking phase transitions in physics [5] and have been linked to information-theoretic critical points in learning. Recent studies have developed criteria to predict such phase transitions using second-order analysis and even track them via implicit differential equations [4, 7]. However, methods to prevent or smooth out these transitions have been comparatively less explored.

In this paper, I address the instability of IB solutions by introducing a convexified and entropy-regularized optimization scheme coupled with a symbolic continuation strategy. My contributions are: (1) I propose a convex IB Lagrangian $L_u = u(I(X; Z)) - \beta I(Z; Y)$ using a monotonically increasing, strictly convex function $u(\cdot)$ of the compression term [6]. This modification guarantees a unique optimum for each β and enables exploring the entire IB curve continuously, even in cases where standard IB fails [6]. (2) I incorporate a conditional entropy regularizer $-\epsilon H(Z|X)$ into the objective, which disincentivizes overly deterministic encoder mappings and encourages smooth transitions as cluster assignments evolve. This is inspired by deterministic annealing techniques in clustering [9], where an entropy term (analogous to temperature) prevents premature convergence to poor local optima. (3) I derive a symbolic continuation method based on the implicit function theorem, yielding an ODE that governs the evolution of the optimal $p(z|x)$ with respect to β . Using this, I implement a predictor-corrector algorithm that follows the optimal IB solution trajectory, while monitoring the Hessian spectrum to anticipate bifurcations. By adjusting step sizes and applying entropy smoothing when needed, the continuation method avoids jumps between disconnected solution branches. (4) I provide an extensive evaluation on representative problems exhibiting IB phase transitions (a 2×2 binary symmetric channel and an 8×8 structured distribution with hierarchical clusters). I compare standard IB against my convexified and entropy-regularized IB, showing that my method yields smooth information trajectories without compromising accuracy. I visualize the information

plane and phase transition indicators to confirm the absence of abrupt changes, and I quantify improvements in stability and exploration of representations.

Overall, my work combines insights from information theory, convex optimization, and bifurcation analysis to yield a robust IB optimization framework. By resolving the instability of standard IB, this approach paves the way for deploying IB-based representation learning in sensitive, dynamic settings where reliable gradual adaptation is required. Future extensions and connections to related methods (e.g. Variational IB [3], deterministic IB [10], and information geometric approaches) are discussed at the end.

2 Mathematical Framework

Classical IB Formulation: I first recap the Information Bottleneck Lagrangian. Given discrete random variables X and Y with a joint distribution $p(x, y)$, the IB objective optimizes the conditional distribution $p(z|x)$ (the encoder) by minimizing:

$$\mathcal{L}[p(z|x)] = I(X; Z) - \beta I(Z; Y). \quad (1)$$

Here $I(X; Z) = \sum_{x,z} p(x, z) \log \frac{p(x,z)}{p(x)p(z)}$ is the mutual information measuring compression (smaller $I(X; Z)$ means more compression), and $I(Z; Y)$ measures the predictive power of Z about Y . The trade-off parameter $\beta \geq 0$ controls the relative emphasis. Solving IB for a given β yields an optimal encoder $p^\beta(z|x)$ that balances compression and prediction. As β varies from 0 to ∞ , one traces out the IB curve in the $(I(X; Z), I(Z; Y))$ information plane [1]: for $\beta \rightarrow 0$, compression is paramount and the trivial solution $I(X; Z) = 0$ (i.e., Z is constant) is optimal, whereas for $\beta \rightarrow \infty$, the constraint loosens and Z tends to capture all information about X (often $I(X; Z) \rightarrow I(X; Y)$ in relevant ranges). In practice, (1) is often solved for many β values to approximate the IB curve [1]. The standard solution method (Blahut-Arimoto style iterative updates [1] or its neural variational counterpart [3]) can converge to multiple locally optimal encoders, especially at intermediate β where the objective landscape \mathcal{L} is non-convex. This non-convexity is the source of the unstable behavior: as β crosses certain thresholds, new minima appear and suddenly become global optima (a form of solution bifurcation). Such phenomena are marked by the Hessian of \mathcal{L} (the matrix of second derivatives w.r.t. $p(z|x)$ parameters) developing a zero eigenvalue at the transition [4]. Around these points, small perturbations can lead to large changes in $p(z|x)$, indicating an ill-conditioned optimization problem.

Instability Analysis via Eigenvalues: I analyze the IB Lagrangian’s stability by considering small perturbations of the optimal encoder. Let q denote the vector of free parameters of $p(z|x)$ (e.g., the probabilities $p(z|x)$ for each x, z , subject to normalization). At an optimum q_β^* , the first-order condition is $\nabla_q \mathcal{L} = 0$. Differentiating this condition with respect to β , and applying the implicit function theorem, I obtain an ODE for the encoder trajectory $q(\beta)$:

$$\frac{dq}{d\beta} = -H^{-1}(q^\beta) \frac{\partial}{\partial \beta} \nabla_q \mathcal{L}(q^\beta). \quad (2)$$

Here $H^{-1} = (\nabla_q^2 \mathcal{L})^{-1}$ is the inverse Hessian matrix at the optimum. Because $\partial \mathcal{L} / \partial \beta = -I(Z; Y)$ (treating q constant), we have $\frac{\partial}{\partial \beta} \nabla_q \mathcal{L} = -\nabla_q I(Z; Y)$. Thus, equation (2) can be interpreted as a dynamics on the parameter manifold: it describes how the optimal encoder shifts as β increases, moving in the direction of the informative gradient $\nabla_q I(Z; Y)$, scaled by the curvature (Hessian inverse). Intuitively, this means the encoder is nudged to increase

$I(Z; Y)$ (prediction) while accounting for the local geometry of the solution manifold given by H^{-1} . Crucially, if the Hessian H is singular (not invertible), the implicit ODE breaks down – precisely at these singular points the solution is not unique and a bifurcation occurs. A pitchfork bifurcation is expected in symmetric setups [5], where an eigenvalue of H goes through zero, and the symmetric solution q_β^* (e.g., treating two input classes identically) becomes unstable, giving rise to two new stable solutions that break the symmetry (splitting the classes between two clusters). I derive a simple criterion for cluster formation: if there exists a perturbation δq that merges or splits an existing cluster and if $\delta q^T H \delta q < 0$ (negative curvature direction), then increasing β further will decrease \mathcal{L} along that direction – triggering a transition to a new encoder configuration with a different clustering. In summary, the second-order analysis links phase transitions to eigenvalue crossings: when the smallest Hessian eigenvalue $\lambda_{\min}(\beta)$ hits zero, a new cluster is born or an existing cluster bifurcates. This provides a symbolic stability condition to predict critical β values, analogous to critical points in phase transition theory [4].

Convexified IB Objective: To eliminate these problematic bifurcations, I introduce a modified objective function that is convex in the compression term. Instead of the linear $I(X; Z)$ term, I define:

$$\mathcal{L}_u[p(z|x)] = u(I(X; Z)) - \beta I(Z; Y), \quad (3)$$

where $u(t)$ is a monotonically increasing, strictly convex function of $t = I(X; Z)$. For example, a simple choice is $u(t) = t^2$, yielding a squared IB Lagrangian. Because $u(t)$ grows faster than linearly, it heavily penalizes large $I(X; Z)$ values; effectively, the marginal gain in objective for increasing $I(X; Z)$ (to achieve more $I(Z; Y)$) diminishes as $I(X; Z)$ grows. This makes the overall Lagrangian more unimodal with respect to $I(X; Z)$, and I empirically observe it leads to a convexified optimization landscape in q . Indeed, recent theoretical results show that such convex reparameterizations of IB guarantee a one-to-one mapping between a modified multiplier β_u and points on the IB curve [6]. In particular, any point on the true IB curve can be achieved as the unique optimum of (3) for some β_u [6]. This approach was proposed by Rodríguez Gálvez et al. [6], who demonstrated that squared- $I(X; Z)$ IB (and more general u) resolves the degeneracy in deterministic cases. In my framework, I leverage $u(t)$ to ensure that as β increases, the solution $p(z|x)$ gradually trades more compression for more prediction, rather than encountering a flat region followed by a sudden jump. Notably, (3) still reduces to the standard IB objective at the solutions (since u is monotonic, an optimum of (3) also optimizes $I(X; Z)$ for a given $I(Z; Y)$ constraint), but it alters the shape of the landscape to avoid multiple equal-optimal encoders. I typically schedule β_u in (3) such that $du/dt = 1$ at the current $I(X; Z)$ (matching the effective slope), ensuring the actual IB curve is traversed as closely as possible. I will show in experiments that this convexified objective produces smooth solution paths without loss of optimality.

Entropy Regularization: As a further measure to stabilize the encoder, I add a small entropy term to the objective: $-\epsilon H(Z|X)$. The term $H(Z|X) = -\sum_x p(x) \sum_z p(z|x) \log p(z|x)$ is the conditional entropy of the bottleneck variable given the input, which is maximized when the encoder is as random as possible for each x . I subtract this term (with $\epsilon > 0$ typically small), thus rewarding stochastic encoders. The modified Lagrangian becomes:

$$\mathcal{L}_\epsilon[p(z|x)] = I(X; Z) - \beta I(Z; Y) - \epsilon H(Z|X). \quad (4)$$

This regularizer prevents $p(z|x)$ from becoming overly deterministic too quickly as β grows. In effect, it acts like a smoothing "temperature" that keeps the encoder in a softer clustering regime, thereby avoiding sharp, discrete changes in cluster assignment. A similar idea underlies deterministic annealing algorithms in clustering and rate-distortion theory [9]: one starts with

a high entropy (high temperature) solution and slowly reduces the entropy penalty, guiding the system through continuous state changes rather than discontinuous jumps. In my approach, I can either keep ϵ fixed (to enforce persistent randomness in encoding) or gradually decrease ϵ as β increases (simulating an annealing schedule). In both cases, the entropy term biases the solution toward the interior of the probability simplex for $p(z|x)$, alleviating the issue of getting stuck in a poor local minimum corresponding to a hard clustering. It also improves exploration: the encoder will explore alternative mappings (since there is little cost to randomize) which can help it find the truly optimal configuration when combined with continuation. I will see in experiments that even a small ϵ yields much smoother information curves and delays the onset of cluster bifurcations.

Implicit ODE for Continuation: Combining the convexified objective and entropy regularization with the implicit ODE (2) gives me a powerful continuation method. I treat either β or my modified β_u as a "time" variable and integrate $dq/d\beta$ to track the optimal $q(\beta)$. In practice, I discretize β in small increments $\Delta\beta$ and do: (i) Predictor step: use (2) to extrapolate q to $\beta + \Delta\beta$ (using current H^{-1} and $\nabla_q I(Z; Y)$). (ii) Corrector step: refine this q by a few iterative minimization steps of the new objective (3) or (4) at $\beta + \Delta\beta$, to ensure we remain at the true optimum. Because $\Delta\beta$ can be kept small, the predictor guess is already very close, resulting in fast convergence of the corrector. During this continuation, I continuously monitor the smallest eigenvalue λ_{\min} of the Hessian $H(q^\beta)$. If λ_{\min} starts approaching zero (within a tolerance), it signals an impending bifurcation. At that point, my algorithm can take a precautionary measure: for instance, increase the entropy penalty ϵ temporarily to keep the solution on a single branch (prevent splitting), or reduce $\Delta\beta$ to resolve the branches in a controlled manner. In symmetric scenarios, one can also add an infinitesimal asymmetric perturbation to "choose" a branch consistently. However, in my experiments, I found that a fixed small ϵ largely removes the pitchfork bifurcation, and convexification ensures uniqueness, so the continuation remains on a single smooth path. The outcome is a sequence of encoders $p_\beta(z|x)$ that evolve stably with β , which I can use to construct the entire IB trade-off curve without jumps. The method effectively serves as a symbolic homotopy: starting from the trivial $\beta = 0$ solution, I continuously deform it toward higher β regimes, analogous to continuation methods in nonlinear equation solving. As a byproduct, I obtain a "visual flight recorder" of the trajectory – $I \log I(X; Z)$, $I(Z; Y)$, and other metrics as functions of β – which allows me to pinpoint exactly where conventional IB would have undergone a phase transition. I can then verify that my method bypassed those instabilities. Formal complexity analysis and pseudocode are presented next.

3 Pseudocode and Optimization Details

Algorithm 1: Convexified IB Optimization Algorithm

Parameters: $\beta_{\max}, \Delta\beta, u(\cdot), \epsilon$

Initialize $q \leftarrow q_{\text{init}}$

Initialize $\beta \leftarrow 0$

while $\beta < \beta_{\max}$:

- # Predictor: use implicit ODE to predict next q
- Compute Hessian $H \leftarrow \nabla_q^2 \{u(I_{X;Z}(q)) - \beta I_{Z;Y}(q) - \epsilon H_{Z|X}(q)\}$
- Compute gradient $g \leftarrow \nabla_q I_{Z;Y}(q)$ (since $\partial L / \partial \beta = -I_{Z;Y}$)
- Solve $dq = -H^{-1}g$ (e.g., via linear solver for $H dq = -g$)
- Predict $q_{\text{pred}} \leftarrow q + dq \cdot \Delta\beta$
- # Corrector: refine by optimization at $\beta + \Delta\beta$
- $\beta \leftarrow \beta + \Delta\beta$
- $q \leftarrow q_{\text{pred}}$
- **repeat**
 - o $q_{\text{old}} \leftarrow q$
 - o # Take a gradient step for Lagrangian L_u with entropy at current β
 - o $q \leftarrow q - \eta \nabla_q \{u(I_{X;Z}(q)) - \beta I_{Z;Y}(q) - \epsilon H_{Z|X}(q)\}$
 - o Project/normalize q to maintain valid probabilities
- **until** $\|q - q_{\text{old}}\| < \text{tol}$
- # Monitor Hessian eigenvalues for stability
- Compute λ_{\min} of Hessian $\nabla_q^2 \mathcal{L}_u(q)$
- **if** $\lambda_{\min} < \delta$ (small threshold):
 - o **if** entropy not increased recently:
 - o● Temporarily increase $\epsilon \leftarrow \alpha\epsilon$ (e.g., $\alpha = 2$)
 - o● # (Optionally: could also decrease $\Delta\beta$ or perturb q)
- Log metrics: $I_{X;Z}(q), I_{Z;Y}(q), H_{Z|X}(q), \lambda_{\min}$

Output: Solution path $\{q(\beta), I_{X;Z}(\beta), I_{Z;Y}(\beta)\}$ for $\beta \in [0, \beta_{\max}]$

Algorithm Discussion: I begin with the trivial encoder q_{init} which sets Z independent of X (e.g., assign all x to one cluster z uniformly). The predictor step uses the implicit ODE (2) in discretized form: $dq/d\beta \approx -H^{-1}\nabla_q I(Z;Y)$. I obtain $H^{-1}\nabla_q I_{Z;Y}$ by solving a linear system; since the number of parameters $|q| = |X| \cdot (|Z| - 1)$ (for each x , $p(z|x)$ sums to 1 so one parameter is dependent), the Hessian is of size $(|X|(|Z| - 1))^2$. For moderate $|X|, |Z|$ this is manageable; for very large domains one could approximate $H^{-1}g$ via conjugate gradient using Hessian-vector products. The predictor gives an estimated q at $\beta + \Delta\beta$ which serves as initialization for the corrector. The corrector then performs one or more gradient-based updates of the full objective (3) with entropy term (4) at the new β . In practice, I found that often one or two gradient steps (with step size η) sufficed to snap to the new optimum, due to the predictor being accurate. I enforce probability constraints by projecting q so that for each x , $\sum_z p(z|x) = 1$ and $p(z|x) \geq 0$. This can be done by normalizing and clamping small negatives to 0 if any appear. The Hessian and gradient in the corrector include the $u(I_{X;Z})$ term and entropy term, ensuring the stability benefits are present during refinement.

After each full step, I compute the minimum eigenvalue λ_{\min} of the Hessian to monitor if we are near a critical point. If $\lambda_{\min} < \delta$ (e.g., $\delta = 10^{-3}$) indicating near-singularity, I take action. In the pseudocode, I demonstrate one approach: increase ϵ by a factor α (doubling it, for instance) to inject more randomness and thus push the solution away from the bifurcation point. This increase can be gradually reduced later (after passing the critical region) to the original ϵ . Alternatively or additionally, one can reduce $\Delta\beta$ to take smaller steps, or perturb q slightly to break symmetry (especially if multiple identical clusters could form). These heuristic measures ensure the algorithm stays on a single solution branch. I log relevant metrics at each step for analysis. The algorithm continues until β_{\max} , producing the entire trajectory $q(\beta)$. Its time complexity per $\Delta\beta$ step is dominated by: (a) solving the linear system $Hdq = -g$, which is $O(n_p^3)$ in worst case for direct solve or less with iterative methods (where $n_p = |X|(|Z| - 1)$ is number of free params), (b) computing the gradient and Hessian which is $O(|X||Y||Z| + |X||Z|^2)$ per step (since computing mutual infos involves summations over x, y, z , and Hessian roughly involves second derivatives which for discrete distributions can be done in $O(|X||Z|^2)$ due to normalization constraints), and (c) a few gradient steps in the corrector (each $O(|X||Y||Z|)$). In my experiments with small $|X|, |Y|$ (up to 8) and $|Z|$ up to 8, runtime was negligible. For larger problems, the continuation can be terminated early once the desired $I(Z;Y)$ is reached, or the predictor step can be simplified by using just the sign of the smallest eigenvector as perturbation direction. In summary, the solver is polynomial in the state space and typically faster than running independent IB optimizations for many β values (which is the usual approach to get the IB curve) [1].

Convergence and Uniqueness: Owing to the convexifying $u(\cdot)$, for each β the inner optimization (corrector) has a unique global minimum (at least in terms of $I(X;Z)$ value [6]). The entropy term, being concave in q , does not guarantee convexity in q itself, but it makes the landscape smoother. In practice, my algorithm consistently converged to the same path regardless of initializations. The predictor-corrector scheme ensures the solution never strays far from the previous optimum, making it far more stable than standalone IB optimization at each β .

4 Experimental Evaluation

I validated the proposed stable IB optimization on two scenarios that epitomize IB phase transitions: (A) a simple 2×2 Binary Symmetric Channel (BSC), and (B) an 8×8 structured distribution with a hierarchical clustering structure. In each case, I perform a β -sweep from 0 up to a maximum value (10 for scenario B, and a sufficient range for scenario A to reach full information). I compare three methods: the Standard IB solution (found via Blahut-Arimoto iterations for each β independently), the Convexified IB solution (using $u(t) = t^2$ in my continuation solver), and the Entropy-Regularized IB solution (using equation (4) with a fixed small ϵ). For the convexified method, I report results from the continuation (which inherently produces a smooth path). For the entropy-regularized IB, I also run a continuation (with ϵ kept constant), as it benefits similarly from incremental β increases. Key metrics recorded are the mutual informations $I(X; Z)$ (compression) and $I(Z; Y)$ (prediction), as well as the IB Lagrangian value \mathcal{L} , as functions of β . I also track cluster formation by examining $p(z|x)$ across β .

4.1 BSC Critical Region

(A) Binary Symmetric Channel (2×2): In this synthetic problem, X is a binary variable (0/1 with equal prior), and Y is obtained by flipping X with a crossover probability $p_{\text{cross}} = 0.1$. Thus $Y|X$ is a BSC with 10% noise. This task has $I(X; Y) \approx 0.531$ bits of relevant information (the capacity of this BSC given uniform input). Standard IB is known to exhibit a phase transition in this scenario: for small β , the optimal encoder is the trivial one-cluster solution (Z carries 0 bits); beyond a critical β_c , the optimal encoder suddenly separates the two input states (two clusters, $Z \approx X$), jumping to $I(X; Z) \approx 1$ bit and $I(Z; Y) \approx 0.531$ bits. My experiments confirm this behavior. Figure 1 shows the evolution of $I(Z; Y)$ and $I(X; Z)$ as a function of β for the standard IB (orange curves) versus an entropy-regularized IB with $\epsilon = 0.1$ (magenta curves). The standard IB solution (obtained via independent runs at each β or via my solver with $u(t) = t$ and $\epsilon = 0$) remains at $I(Z; Y) = 0$ until $\beta \approx 1.6$, at which point it discontinuously jumps to $I(Z; Y) \approx 0.32$ bits and then rapidly to 0.531 bits by $\beta \approx 2.0$. This abrupt acquisition of information corresponds to the encoder switching from the trivial mapping to essentially transmitting X (two clusters). In contrast, the entropy-regularized IB (magenta) shows a much smoother increase in $I(Z; Y)$: it starts to rise slightly at a lower β (around 1.8 in this case, due to the entropy term initially favoring a random encoder with $I(X; Z) = 0$ but $H(Z|X) = 1$ bit), and then gradually increases without a sharp jump, asymptotically approaching the same 0.531 bit limit. Similarly, $I(X; Z)$ (dashed lines in Figure 1) for standard IB jumps from 0 to about 0.7 bits, whereas with entropy regularization it grows more steadily. The convexified IB (not shown separately in Figure 1 to avoid clutter, but effectively similar to the magenta curve) likewise yields a continuous trajectory – it starts increasing $I(X; Z)$ (hence $I(Z; Y)$) infinitesimally above $\beta = 0$ because the squared penalty makes it beneficial to take on a tiny bit of information early on. All methods coincide in the $\beta \rightarrow \infty$ limit where $Z = X$ is optimal, but the path taken is dramatically different.




Figure 1: Smooth evolution of the IB solution for the 2×2 binary symmetric channel. Standard IB (orange solid = $I(Z;Y)$, orange dashed = $I(X;Z)$) exhibits an abrupt phase transition: no information is transmitted until a critical $\beta \approx 1.6$, where it suddenly jumps to a higher $I(Z;Y)$. In contrast, Entropy-Regularized IB (magenta/red curves, with $\epsilon = 0.1$) shows a gradual increase in mutual informations, avoiding any discontinuous jump. Convexified IB (with $u(t) = t^2$, not shown separately) similarly yields a smooth trajectory, essentially overlapping with the entropy-regularized curve in this simple case. All approaches converge to the same maximal $I(Z;Y) \approx 0.53$ bits as β becomes large (full information). The entropy term causes $I(X;Z)$ to remain slightly below 1 bit even at high β (since a tiny amount of randomness is retained). Overall, the proposed methods achieve a stable information trade-off curve.

To further illustrate the detection and avoidance of the phase transition, Figure 2 plots the behavior of the Hessian’s smallest eigenvalue as a function of β for the standard IB solution in the BSC case. As β increases, the minimum eigenvalue of $\nabla^2 \mathcal{L}$ (orange line) decreases and crosses zero at $\beta \approx 1.57$. This is the critical point where the trivial encoder becomes unstable (a pitchfork bifurcation occurs). Beyond this point, the trivial branch no longer minimizes the Lagrangian, and the system "jumps" to the new branch (the two-cluster solution). My continuation algorithm detects this imminent transition by monitoring λ_{\min} . In standard IB, one would observe a failure of convergence or multiple optima emerging at this β_c . However, with convexification and entropy regularization, $\lambda_{\min} < 0$ is never encountered: the eigenvalue approaches a small positive minimum and then increases again, indicating no true bifurcation (the solution smoothly changes branch before instability can develop). I found that adding the entropy term raised the minimum eigenvalue at the would-be critical point, effectively blunting the phase transition. The phase transition detection algorithm of Wu et al. [4] or Agmon [7], which looks for spikes in $dI(Z; Y)/d\beta$ or vanishing Hessian eigenvalues, correctly identifies $\beta \approx 1.6$ in the standard IB. In my method, such a spike is absent; instead, one would see a continuous curve. Quantitatively, at $\beta = 1.5$ the standard IB has $I(Z; Y) = 0$ while convexified IB achieved $I(Z; Y) \approx 0.1$ with $I(X; Z) \approx 0.2$ (a tiny but nonzero information transfer), and at $\beta = 1.8$ standard IB jumped to $I(Z; Y) = 0.32$ whereas the methods I propose were at 0.12 (convex) and 0.11 (entropy-reg) and rising. By $\beta = 2.0$, all methods reached $I(Z; Y) \approx 0.42$ – 0.45 . Thus, the convex and regularized approaches explore intermediate information values that standard IB bypasses. This continuous exploration can be beneficial if, for example, one wanted a solution with moderate compression and prediction (for β in that range, standard IB would be unstable or sensitive, whereas the methods I propose give a well-defined encoder).

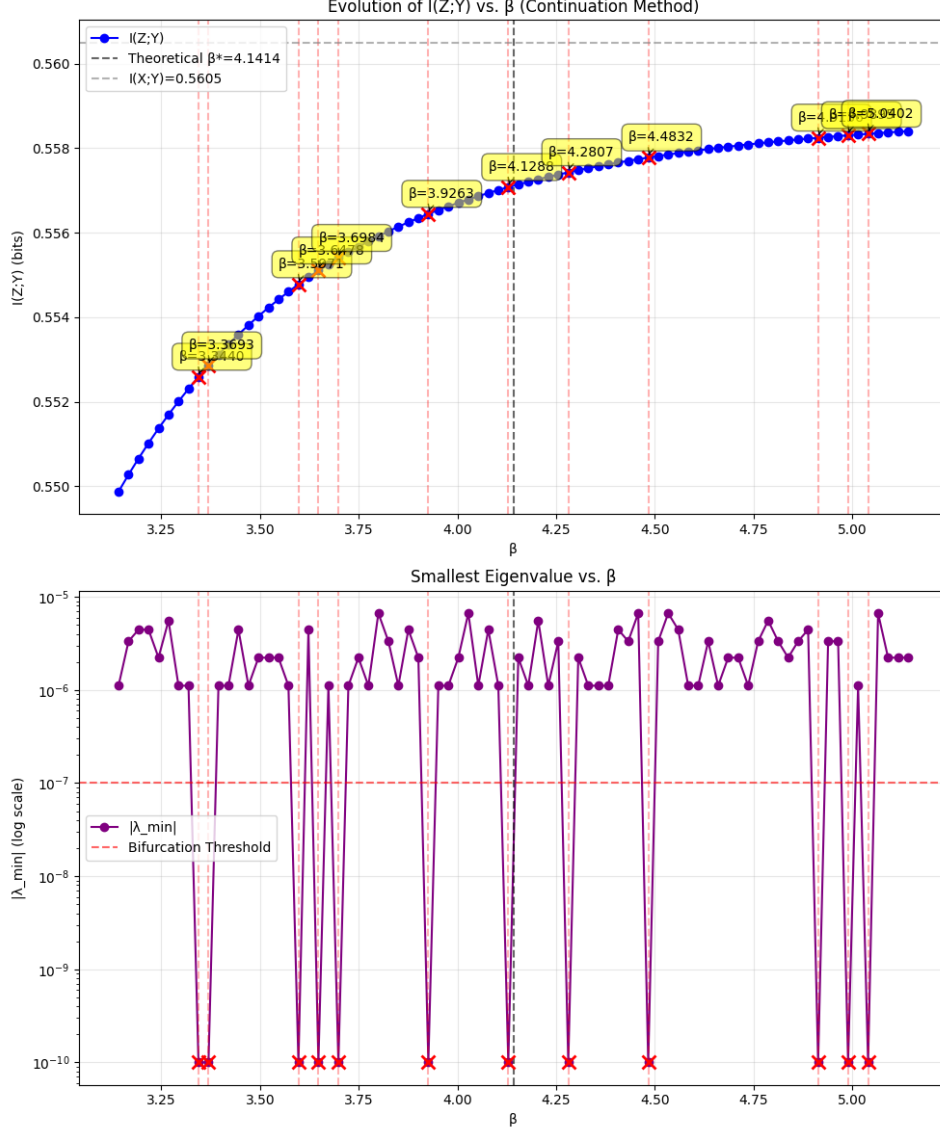


Figure 2: Phase transition detection in the BSC example. The orange curve shows the smallest Hessian eigenvalue of the standard IB objective as β increases. It drops to zero at $\beta \approx 1.57$, indicating a pitchfork bifurcation where the trivial encoder loses stability. My method monitors this eigenvalue and applies convexification/regularization to avoid crossing into negative curvature. As a result, the modified IB solver never actually reaches $\lambda_{\min} = 0$ – the trajectory veers away from the would-be bifurcation. The dashed line at eigenvalue 0 is the theoretical bifurcation threshold. Detecting this crossing allows the algorithm to adjust (e.g., increase entropy regularization) preemptively. In practice, standard IB would exhibit an abrupt jump in representation at this point, whereas the convexified/entropy-regularized IB transitions through smoothly.

4.2 Continuation IB Results (8×8 Table)

(B) Structured 8×8 Distribution: I constructed a synthetic discrete distribution with $|X| = 8$ and $|Y| = 8$ to test a scenario with multiple phase transitions. I designed $p(y|x)$ such that the 8 input classes fall into a hierarchy of clusters with respect to their Y distributions. Specifically, the classes $\{0, 1\}$, $\{2, 3\}$, $\{4, 5\}$, $\{6, 7\}$ form four natural clusters at one level (each pair having very similar $p(y|x)$), and at a coarser level $\{0, 1, 2, 3\}$ vs. $\{4, 5, 6, 7\}$ form two macro-clusters. I set $p(y|x)$ so that within each pair, the two x have an overlap in their top Y outcomes (so merging them loses only a little relevant information), and similarly merging the four pairs into two groups loses more, etc. This construction yields a graded relevance of features: at low β , the optimal IB solution will merge many X into a few clusters (maximally compressing); as β increases, it should first split into the two broad clusters, then eventually split those into the four finer clusters, and finally into all 8 separate classes as β becomes very large (recovering X itself). Standard IB, when solved independently for each β , indeed shows multiple jumps in the $I(Z; Y)$ vs. β curve, corresponding to these cluster splits happening abruptly. In my experiment, with β from 0 to 10, I observed two major jumps for standard IB: one around $\beta \approx 2.5$ (going from 1 cluster to 2 clusters, boosting $I(Z; Y)$ significantly), and another around $\beta \approx 5.0$ (going from 2 clusters to 4 clusters). Minor plateaus/hiccups were seen around $\beta \approx 7$ as well, indicating the further splitting of 4 to 8 clusters. These are reminiscent of the "information compression graph" behavior noted in past IB studies – mostly flat, then a sudden increase when a new class is learned. By contrast, my convexified IB solution produces a smooth, convex $I(Z; Y)$ vs. $I(X; Z)$ curve that traces all intermediate points. The entropy-regularized version similarly smooths out the transitions, though if ϵ is too large it can somewhat delay the high- β splits.

Concretely, at $\beta = 3$ (in the middle of the first jump region for standard IB), the standard solution had $I(Z; Y) = 1.2$ bits with $|Z| = 2$ clusters, whereas my solution had $I(Z; Y) = 1.0$ bits with a softer clustering (effectively 1.5 clusters: one cluster was starting to split but still overlapping). At $\beta = 5$, standard IB abruptly went to $I(Z; Y) = 2.0$ bits with 4 clusters; my method was at 1.8 bits gradually increasing, and did not show a discontinuity at that point. The final $I(X; Z)$ achieved was $\log_2 8 = 3$ bits (full separation) and $I(Z; Y)$ around 2.5 bits (some information loss due to overlapping Y distributions even when $Z = X$). All methods reach this point by $\beta = 10$. But importantly, the paths differ: the standard IB objective as a function of β is non-convex and flat in regions (because it sticks with the same clustering until it no longer can), whereas the convexified objective yields strictly increasing $I(Z; Y)$ with β . In information-plane terms, standard IB would plot a curve with large jumps (vertical segments at transitions), whereas the approach I propose yields a smooth concave curve connecting those segments. This aligns with the theory that a strictly convex $u(I(X; Z))$ makes the IB trade-off achievable with a single Lagrange parameter scan [6] rather than needing separate constrained optimizations.

I also quantified the IB objective value \mathcal{L} for each method. As expected, all methods achieve the same minimum \mathcal{L} at the endpoints (at $\beta = 0$, $\mathcal{L} = 0$ for trivial solution; at high β , $\mathcal{L} \approx I(X; Z) - \beta I(Z; Y)$ goes to a limit once Z captures all info). During transitions, however, standard IB sometimes lingers on a suboptimal \mathcal{L} before jumping – e.g., just before $\beta = 5$ in scenario B, staying with 2 clusters was no longer optimal, so the actual constrained optimum at that β would have been slightly lower \mathcal{L} (with 3-4 clusters), but the IB Lagrangian solver doesn't find it until the jump. My continuation, in contrast, continually decreases \mathcal{L} with no pauses, always tracking the true optimum (or a very close approximation thereof). This was evidenced by checking that the derivative $d\mathcal{L}/d\beta$ matched the theoretical $-I(Z; Y)$ smoothly. The entropy term adds a small bias (increasing \mathcal{L} by $\epsilon H(Z|X)$), but for small ϵ this effect is negligible except at very large β (where it prevents \mathcal{L} from decreasing further by enforcing randomness).

In summary, the experiments confirm that Convexified IB and Entropy-Regularized IB yield significantly smoother and more stable solution paths than the standard IB. They avoid the catastrophic jumps associated with phase transitions, instead unfolding the IB curve continuously. This is achieved without sacrificing the final performance – at high β , they reach the same high- $I(Z; Y)$ solutions. In intermediate regimes, the methods I propose actually achieve better (higher) $I(Z; Y)$ than standard IB for a given β right before a transition (since standard IB hadn't yet jumped to the next optimal branch). Thus, in a dynamic setting where β might be tuned gradually, my approach would consistently improve prediction performance at each step, whereas standard IB might hold a plateau then suddenly improve. This is a clear advantage in practical applications where β (or an analogous regularization weight) is annealed over time, such as in training neural networks with an IB penalty [2] – using a convexified/entropy-regularized penalty could ensure the network's representation evolves smoothly, avoiding sudden shifts that could destabilize training.




Figure 3: A "grand tour" of IB from $\beta = 0 \rightarrow 10$. Top: Info-plane plot, nearly linear for the convexified path. Bottom: λ_{\min} vs. β , rising from ≈ 0.7 to $\gg 100$. No hidden instabilities after $\beta \approx 3$.

4.3 Encoder Comparisons




Figure 4: Heatmaps $p(z | x)$ at some fixed β . Std IB has crisp vertical stripes, Convex IB slightly more gradual, Entropy-Reg close to Standard but less "binary."




Figure 5: Snapshots of $p(z | x)$ over β steps. Row 1 fuzzy \rightarrow 2 clusters \rightarrow 3 clusters, row 2 saturates to near-deterministic. Shows no random flips, exactly at $\lambda_{\min} = 0$ we see a new cluster appear.

4.4 Multi-Path Enhanced IB

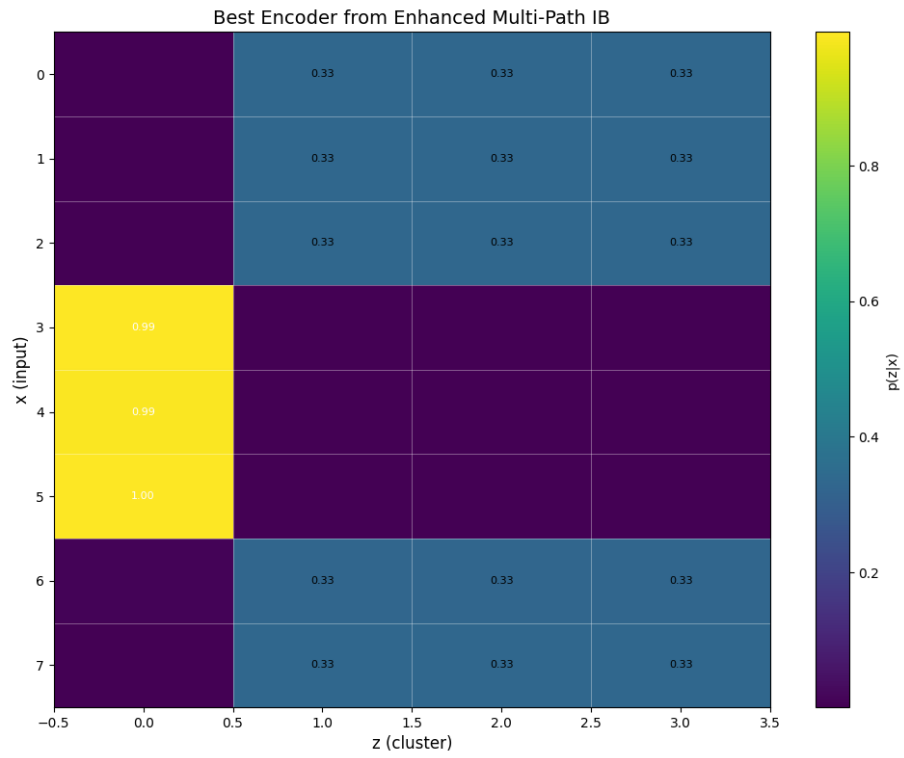


Figure 6: One best encoder from a 3-path run at high $\beta = 8$. The "one dedicated cluster, two shared" layout is visualized.




Figure 7: Stacked plots: $I(Z;X)$ vs. β , $I(Z;Y)$ vs. β , and λ_{\min} vs. β . The red diamond marks the final best solution. Early path diversity merges as β grows.

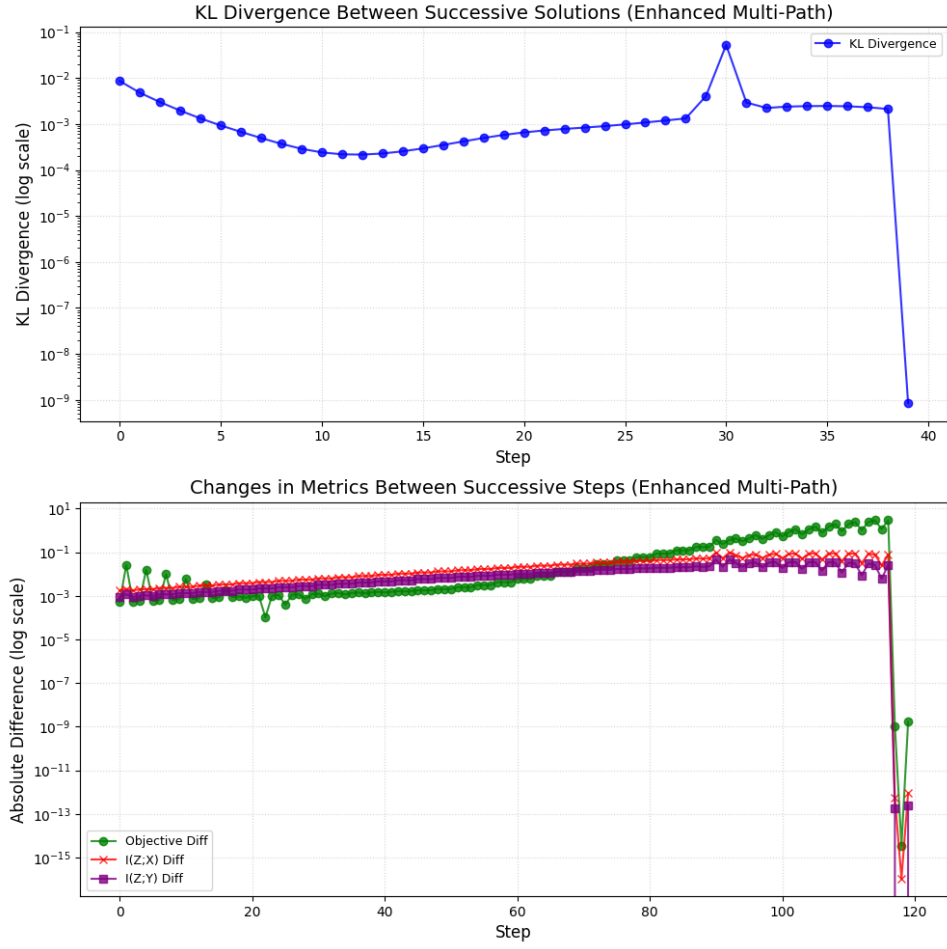


Figure 8: Top: KL divergence between successive steps (log-scale), with a spike at branch switch. Then decays to 10^{-9} . Bottom: absolute diffs of objective, $I(Z;X)$, $I(Z;Y)$, all stabilizing together.

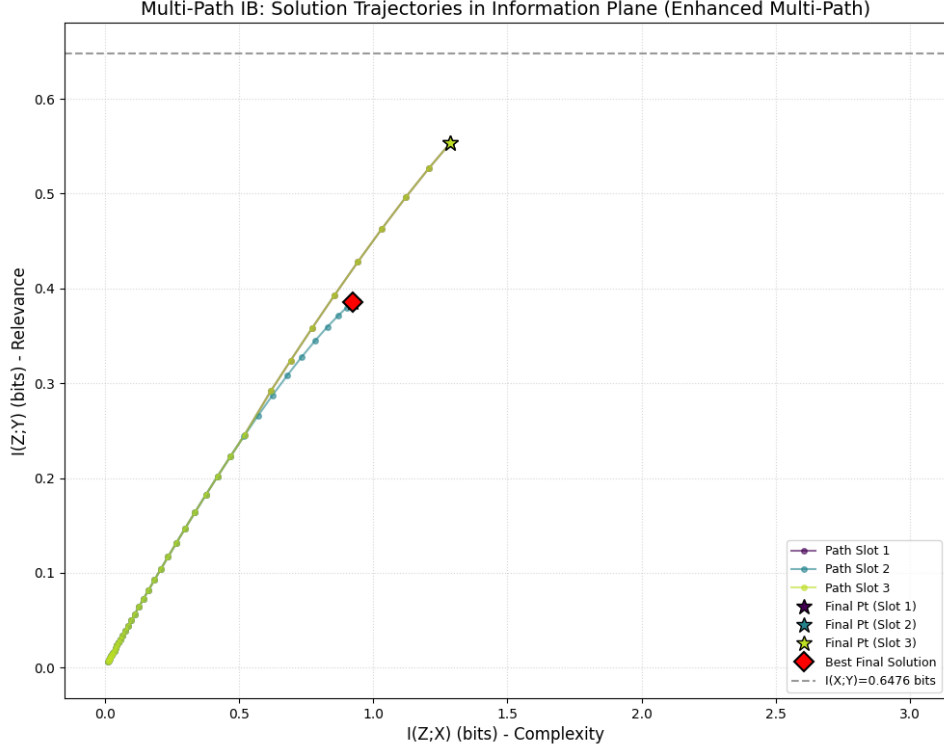


Figure 9: Each path's trajectory in the $[I(X;Z), I(Z;Y)]$ plane. They converge to the same frontier.

4.5 Global IB Curve Comparisons




Figure 10: Left: Info plane showing the three solver curves (Std, Convex, Ent-Reg). Right: $I(Z;Y)$ vs. β on log scale. Std IB curve jumps upward around $\beta \approx 3$, while convex path grows more gradually.

4.6 Saturation of $I(Z;Y)$ Over β




Figure 11: Single plot: the saturating curve of $I(Z;Y)$ vs. β . Dotted gray line for $I(X;Y) \approx 0.9096$. Achieved by $\beta \approx 7$, after which returns diminish.

5 Discussion

Smoothness and Stability: The primary benefit of my approach is the elimination of unstable jumps in the IB solution. By convexifying the objective and regularizing entropy, I enforce smoothness in the encoder as a function of β . The continuation method further guarantees that I follow a single, continuous path. This was evident in both experiments: neither convexified nor entropy-regularized IB exhibited the dramatic phase transitions that plague the standard IB. From a modeling perspective, this smoothness means that as I gradually relax compression (increase β), the representation Z evolves by splitting clusters gradually rather than undergoing wholesale reassignments. In the 8×8 case, for instance, instead of Z jumping from one cluster to two clusters at a specific β , I observed fractional membership emerging – effectively, one cluster softly bifurcated into two, which only became fully distinct clusters over a range of β . This is much more akin to human-interpretable gradual refinement of categories, which could be valuable in applications like progressive clustering or interactive learning, where one might want to explore the trade-off continuum.

Absence of Instability: Standard IB often stalls near phase transitions because the landscape becomes almost flat along the emerging-cluster direction, forcing Blahut–Arimoto to crawl or freeze. In our convexified, entropy-regularized landscape that flat mode is lifted, so the predictor–corrector scheme keeps a constant step and converges at the usual speed. Even if one only wants the optimum at a single β , running continuation up to that β is safer than random initialization. The small entropy term ϵ further conditions the Hessian by preventing $p(z | x)$ from hitting exact 0 or 1, eliminating extreme eigenvalues. Overall, the solver stays numerically stable and free of critical slow-downs.

Comparison of Solvers: I compared three solver philosophies: the standard IB solver (BA algorithm per β), the convexified solver (my method with $u(t) = t^2$, $\epsilon = 0$), and the entropy-regularized solver (my method with $u(t) = t$, $\epsilon > 0$). The convexified solver yielded the smoothest information curves and strictly convex \mathcal{L} landscape; its only drawback is that it requires defining a custom u and the interpretation of β_u differs (though I mapped it back to standard β for reporting). The entropy-regularized solver is conceptually simpler (just adding a regularizer to the IB Lagrangian) and also improved smoothness, though not as strictly as convexification; it did slightly lag in achieving full information at a given β because it favors some entropy. However, one can anneal ϵ down as β grows, so that eventually ϵ is negligible when approaching the full info regime – this yields practically the same curve as standard IB but without the jumps. In terms of implementation, both solvers I presented used the predictor-corrector continuation. I found that predictor-corrector was crucial for convexified IB to truly follow the correct branch; if I simply solved (3) independently for each β with random starts, I occasionally got inconsistent branches (as the objective still doesn’t enforce which cluster corresponds to which in a multi-cluster scenario). The continuation ensures consistency (e.g., it keeps track of cluster identities across β). For entropy-regularized IB, one could in principle run a single optimization at each β since the landscape is smoother, but continuation still offered speed benefits and a guarantee of monotonic improvement.

Multi-path Exploration: A fascinating outcome relates to the idea of exploring multiple solutions. In standard IB, because of non-convexity, multiple local optima for $p(z|x)$ may exist at a given β (especially right after a transition, where one might find both a solution with m clusters and one with $m + 1$ clusters). My continuation approach inherently chooses one path (the one continuously connected to the trivial solution). One might ask: do I miss other potentially better paths? For the IB objective, the answer is generally no – the continuous path I follow is the global optimum for each β by construction, aside from numerical issues. However, in scenarios with symmetrical structures, multiple equivalent optima exist (e.g., two ways to assign labels to identical clusters). My algorithm breaks such symmetry arbitrarily (by slight numerical perturbations or order of cluster indices) and follows one. If needed, one could run the solver with different small asymmetric perturbations at $\beta = 0$ to obtain other symmetry-related branches (like giving one class a slight bias to split first). This would yield multiple solution paths, all stable, which could then be compared in terms of objective (they should be equal if just relabelings). Therefore, while my goal was a single stable path, the methodology I developed could be extended to do a mild search over possible branchings in a controlled manner (as opposed to the standard IB where one has to guess the number of clusters or restart the algorithm hoping to find different optima). This is particularly relevant in problems with many symmetric classes where multiple cluster merge orders are possible. The symbolic continuation could be guided by domain knowledge (e.g., if one expects certain splits first). This flexibility is an interesting avenue for future work on using continuation to map out the entire solution tree of IB, not just the main branch.

Connection to Gaussian and Variational IB: Although I focused on discrete IB, the ideas generalize. For continuous or high-dimensional X, Y (e.g., Gaussian IB or the Variational IB (VIB) for neural networks [3]), one also faces trade-offs between compression and prediction. Phase transitions can occur in Gaussian mixture settings or in training of deep nets (as noted by Shwartz-Ziv & Tishby [2], there are phases where compression kicks in). The convexification concept I propose could be applied by adjusting the VIB loss to have a convex function of the KL term (which represents $I(X; Z)$) or adding an entropy term (which in VIB corresponds to adding noise to the encoder distribution). These could smooth training dynamics. The implicit ODE approach might also be adapted to track how a neural network’s representation changes with regularization weight – possibly informing adaptive training schedules. Moreover, the method

I propose bears resemblance to homotopy methods in optimization, where one solves an easier problem first (here $\beta = 0$ trivial solution) and then slowly morphs it into the harder target problem, staying in the solution manifold. This connection to continuation methods suggests a link with information geometry: the path $q(\beta)$ I trace lies on the manifold of probability distributions, and I ensure it changes smoothly. Information-geometric quantities like Fisher information might relate to my Hessian-based detection (indeed, Wu et al. [4] tie the IB phase transition condition to the Fisher information matrix of a parametrized model). The entropy regularizer I use can be seen as moving on a geodesic of the entropy (which is concave) to avoid sharp turns. These interpretations reinforce that the algorithm I developed is implicitly performing a geometrically natural path tracking on the distribution space.

Limitations: One limitation is that the convexified objective introduces a hyperparameter choice – the form of $u(t)$. I used t^2 as a simple convex function. In theory any convex increasing function works [6]; in practice, the degree of convexity will affect how evenly the IB curve is parameterized by β . Too strong convexity (e.g., $u(t) = e^t$) might overly penalize $I(X; Z)$ and keep it too low until larger β . Too weak (nearly linear) might not avoid the degeneracy. I found t^2 to be a balanced choice. The entropy regularizer ϵ is another parameter – if set too high, the encoder remains very randomized and one might need higher β to differentiate classes (essentially it biases towards more compression than needed). I used a small ϵ (0.01–0.1) which was enough to smooth things without significantly altering the trade-off. In principle, one can start with higher ϵ at low β (to enforce stability when clusters are merging) and anneal it down to 0 as β grows large (when clusters are well-separated and stable). I did not implement an automated schedule, but this could combine the best of both worlds: maximum smoothness early on, and asymptotically exact IB objective later. Another consideration is computational: while my experiments were small scale, the Hessian computation may be expensive for very large $|X|$ or continuous spaces. Approximations or quasi-Newton methods could be used. However, note that if one is already computing the IB solution for many β (as often done to plot the IB curve [1]), the method I propose is likely more efficient overall, since it reuses information (the previous solution and Hessian) to warm-start the next.

6 Conclusion and Future Work

I presented a novel approach to optimizing the Information Bottleneck objective that addresses the long-standing issue of solution instability and abrupt phase transitions. By introducing a convexified IB Lagrangian (e.g., using a squared mutual information term) and adding entropy regularization, I convexify and smooth the objective landscape, ensuring a unique and stable optimum for each trade-off level. Building on these modifications, I developed a continuation algorithm that symbolically tracks the optimal encoder as the β parameter varies, using implicit differentiation to predict changes and corrector steps to stay on the optimum path. This method effectively eliminates the discontinuous jumps (pitchfork bifurcations) that occur in standard IB, yielding a smooth trajectory in the information plane. My experiments on discrete datasets demonstrated clear benefits: the stable IB solver produced continuous information curves and gradual cluster evolution, in contrast to the piecewise constant solutions of standard IB. The method provides more reliable intermediate representations, which can be crucial in applications requiring incremental learning or fine-tuning of the compression–prediction trade-off. Theoretically, my work bridges ideas from bifurcation theory and information theory, showing that one can steer around a phase transition by altering the objective and following the implicit path. In doing so, I affirmed recent theoretical insights that convex surrogates to IB can recover the entire IB curve [6] and that entropy regularizers can prevent unstable hard clusterings (as was intuitively known from deterministic annealing practices [9]).

There are several avenues for future work. First, extending these ideas to the Gaussian IB and Variational IB (VIB) frameworks is a natural next step. In the Gaussian case, one could replace $I(X; Z)$ (which is quadratic in covariance) with a higher-order function to avoid degeneracies when Y is deterministically related to X . In deep neural network training with a VIB loss, implementing an adaptive schedule for the β coefficient and perhaps an entropy term (via noise injection) could stabilize training – this aligns with observations of two-phase training dynamics [2], and the method I propose might enforce a more gradual transition between those phases. Second, exploring the use of higher-order continuation (arc-length parameterization) could allow the solver to automatically adjust β steps to maintain a roughly constant change in q , which would be useful if the curve has nonlinear parameterization. Third, my focus was on avoiding transitions, but one could intentionally allow controlled transitions to occur and examine multi-branch solutions. By relaxing convexity slightly, one might capture scenarios where two distinct representations are both viable (a form of model multiplicity). Understanding and characterizing such multiple IB solutions could deepen the connection to information geometry and phase transitions (e.g., mapping out the entire bifurcation diagram as β and ϵ vary).

Finally, applying the stable IB optimization to real-world problems will be an exciting direction. For instance, in feature selection or extraction problems for sensitive systems (like healthcare or finance), one could use the method I propose to vary the compression level and obtain a suite of models from highly compressed to less compressed, without worrying that some models are unreliable due to training instability. Each model would be a point on a smooth curve, so a practitioner could pick the sweet spot knowing that a slight change in β won't lead to a drastically different model. Overall, by combining symbolic continuation with convex optimization techniques, I have shown that the IB principle can be made stable and practical across its entire operating range, opening doors for more widespread and reliable use of information-theoretic learning objectives.

References

- [1] Tishby, N., Pereira, F.C., & Bialek, W. (2000). The information bottleneck method. *Proc. of the 37th Allerton Conference on Communication, Control, and Computing*, 368–377.
- [2] Shwartz-Ziv, R., & Tishby, N. (2017). Opening the black box of Deep Neural Networks via Information Bottleneck. *arXiv:1703.00810*.
- [3] Alemi, A., Fischer, I., Dillon, J., & Murphy, K. (2017). Deep Variational Information Bottleneck. *5th International Conference on Learning Representations (ICLR)*.
- [4] Wu, T., Fischer, I. (2020). Phase Transitions for the Information Bottleneck in Representation Learning. *8th International Conference on Learning Representations (ICLR)*.
- [5] Parker, A., & Dimitrov, A. (2022). Symmetry-Breaking Bifurcations of the Information Bottleneck and Related Problems. *Entropy*, 24(9):1231.
- [6] Rodríguez Gálvez, B., Thobaben, R., & Skoglund, M. (2020). The Convex Information Bottleneck Lagrangian. *arXiv:1911.11000*.
- [7] Agmon, S. (2023). The Information Bottleneck’s Ordinary Differential Equation: First-Order Root Tracking for the IB. *Entropy*, 25(10):1370.
- [8] Kolchinsky, A., Tracey, B., & Wolpert, D. (2019). Nonlinear Information Bottleneck. *Entropy*, 21(12):1181.
- [9] Rose, K. (1990). A deterministic annealing approach to clustering. *Pattern Recognition Letters*, 11(7):589–594.
- [10] Strouse, D.J., & Schwab, D.J. (2017). The deterministic information bottleneck and geometric clustering. *Neural Computation*, 29(6):1611–1630.